

PyTorch Workshop

João Oliveira Parracho

18 December 2024

Table of Contents

- What is PyTorch?
- Tensors
- Loading Datasets
- Building a model
- Training loop
- Inference

What is PyTorch?



PyTorch

- **Open-source** machine learning framework developed by **Meta AI** (Facebook).
- Widely used for **signal processing tasks** such as **image denoising, audio enhancement, and compression**.
- User friendly for researchers and developers to build deep-learning models
- Strong **GPU support**
- **Pytorch website**: <https://pytorch.org/>




PyTorch installation

PyTorch Build	Stable (2.5.1)			Preview (Nightly)	
Your OS	Linux		Mac		Windows
Package	Conda	Pip		LibTorch	Source
Language	Python			C++ / Java	
Compute Platform	CUDA 11.8	CUDA 12.1	CUDA 12.4	ROCm 6.2	CPU
Run this Command:	pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu118				


Previous versions of PyTorch >



What are tensors?



Tensors: Scalar




Numpy:

```
1 import numpy as np
2
3 x = 1
4
```

Torch:

```
1 import torch
2
3 x = torch.Tensor(1)
4
```

Tensors: 1D Array




Numpy:

```
1 import numpy as np
2
3 x = np.array([1.0, 2.0, 3.0])
4
```

Torch:

```
1 import torch
2
3 x = torch.Tensor([1.0, 2.0, 3.0])
4
```


Tensors: 2D Array




Numpy:

```
1 import numpy as np
2
3 x = np.array([[1.0, 2.0, 3.0],
4               [4.0, 5.0, 6.0]])
5
```

Torch:

```
1 import torch
2
3 x = torch.Tensor([[1.0, 2.0, 3.0],
4                   [4.0, 5.0, 6.0]])
5
```

Tensors: 3D Array



Numpy:

```

1  import numpy as np
2
3  x = np.array([[[1.0,2.0],[3.0,4.0]],
4                [[5.0,6.0],[7.0,8.0]]])
5

```


Torch:

```


1  import torch
2
3  x = torch.Tensor([[[1.0,2.0],[3.0,4.0]],
4                    [[5.0,6.0],[7.0,8.0]]])
5

```

Tensors: 3D Array



Tensors: 4D Array



Numpy:

```

1  import numpy as np
2
3  x = np.array([[[[1,2],[3,4]],[[5,6],[7,8]]],
4               [[[1,2],[3,4]],[[5,6],[7,8]]]])
5

```

Torch:

```

1  import torch
2
3  x=torch.Tensor([[[[1,2],[3,4]],[[5,6],[7,8]]],
4                 [[[1,2],[3,4]],[[5,6],[7,8]]]])
5

```

Tensors: Initialisation

```
1  import numpy as np
2  import torch
3
4  # Initialising with zeros
5  x = np.zeros((2,2))
6  y = torch.zeros((2,2))
7
8  # Initialising with ones
9  x = np.ones((2,2))
10 y = torch.ones((2,2))
11
12 # Initialising randomly
13 x = np.random.rand(2,2)
14 y = torch.rand((2,2))
15
```

Tensors: Attributes

```
1
2  import torch
3
4  x = torch.rand((4,3,20,20))
5
6  print(x.shape) # torch.Size([4, 3, 20, 20])
7  print(x.type()) # torch.FloatTensor
8  print(x.device) # cpu
9
```

Tensors: From numpy to PyTorch, and CUDA

- Convert numpy array to Pytorch tensor:
 - **torch.from_numpy()**
- Convert Pytorch tensor to numpy array:
 - **y.numpy()**
- Change Pytorch tensor device:
 - **y.to()**

```

1  import torch
2  import numpy as np
3
4  x = np.random.rand(4,3,20,20)
5  # Convert numpy array to PyTorch tensor
6  y = torch.from_numpy(x) # Returns a cpu
   tensor
7  # Convert PyTorch tensor to numpy array
8  z = y.numpy()
9
10 if torch.cuda.is_available():
11     device = "cuda"
12 else:
13     device = "cpu"
14 # Load the tensor to the cpu or cuda (
   gpu)
15 y = y.to(device)
16

```

Tensors: Operations

```
1  import torch
2
3  x = torch.rand((4,3,20,20))
4
5  # Indexing
6  x_1 = x[:, :, 0:10, :] # torch.Size ([4, 3, 10, 20])
7  x_2 = x[:, :, 10: , :] # torch.Size ([4, 3, 10, 20])
8  x_3 = x[:, :, :, 0:10] # torch.Size ([4, 3, 20, 10])
9  x_4 = x[:, :, :, 10:] # torch.Size ([4, 3, 20, 10])
10
11 # Concatenate tensors
12 y_1 = torch.cat((x_1,x_2),dim=2) # torch.Size ([4, 3, 20, 20])
13 y_2 = torch.cat((x_3,x_4),dim=3) # torch.Size ([4, 3, 20, 20])
14
15 # Swap tensors dimension
16 z = x.permute(3,0,2,1) # # torch.Size ([20, 4, 20, 3])
17
```




ng

CIFAR-10 Dataset

- 60000 32×32 RGB images
- It has **10 classes**, with 6000 images per class.
- **Training images:** 50000
- **Test images:** 10000

airplane




automobile



bird




cat



deer




dog



frog



horse



ship



truck



Dataset: training, validation, and test set

- **Data Division:** Split data into three sets — **training, validation, and test**
- **Training Phase:** Train the model using the training set to adjust the initial parameters.
- **Validation Phase:** Evaluate the model in the validation set and **fine-tune the initial parameters** based on these results.
 - Usually validation set is around 20%,25% of the total training set images
- **Inference:** Evaluate the trained model on the test set to verify its real-world performance.

Create a custom Dataset class

- The custom dataset class must **inherit** the abstract **torch.utils.data.Dataset** class
- **It must implement 3 methods:**
 - `__init__`
 - A constructor method used to initialise the network's parameters.
 - `__len__`
 - Returns size of the dataset
 - `__getitem__`
 - Enables indexing so that `dataset[i]` retrieves the i^{th} sample.

```
1 from torch.utils.data import Dataset
2 from torchvision import transforms
3 from torchvision import datasets
4
5 class Custom_Dataset(Dataset):
6     def __init__(self, split="train"):
7         is_train = True if split == "train" or split == "val" else False
8
9         transform = transforms.ToTensor()
10        self.dataset = datasets.CIFAR10('data', train=True,
11                                       download=True,
12                                       transform=transform)
13
14        # lets use 20% of the training dataset to validation
15        if split == "train":
16            self.indices = self.indices[0:int(len(self.dataset)*0.8)]
17        elif split == "val":
18            self.indices = self.indices[int(len(self.dataset)*0.8):]
19        elif split == "test":
20            # self.dataset.targets: returns the image classes (0,1,...,9)
21            self.indices = np.array(self.dataset.targets)
```

```
1
2  def __getitem__(self, idx):
3      image = self.dataset[self.indices[idx]][0]
4      label = self.dataset[self.indices[idx]][1]
5      return image, label
6
7  def __len__(self):
8      return len(self.indices)
9
```

How to load images from the custom dataset?

- **torch.utils.data.DataLoader** :
Provides an iterable over the given dataset.
- **dataset**: Dataset from which to load the data
- **batch_size**: How many images loaded per batch
- **shuffle**: Enable data reshuffled at every epoch
- **num_workers**: Number subprocesses to use for data loading.

```
1  from torch.utils.data import DataLoader
2
3  train_set = Custom_Dataset(split="train")
4
5
6  trainloader=DataLoader(train_set ,
7                          batch_size=4,
8                          shuffle=False ,
9                          num_workers=2)
10
11
```

How to load images from the custom dataset?

```
1  from torch.utils.data import DataLoader
2
3  train_set = Custom_Dataset(split="train")
4  val_set = Custom_Dataset(split="val")
5  test_set = Custom_Dataset(split="test")
6
7  trainloader = DataLoader(train_set, batch_size=4, shuffle=False, num_workers=2)
8  valloader = DataLoader(val_set, batch_size=4, shuffle=False, num_workers=2)
9  testloader = DataLoader(test_set, batch_size=4, shuffle=False, num_workers=2)
10
11 for i, data in enumerate(trainloader):
12     inputs, labels = data
13
14 for i, data in enumerate(valloader):
15     inputs, labels = data
16
17 for i, data in enumerate(testloader):
18     inputs, labels = data
19
```

Create Custom Balanced dataset

- The CIFAR-10 has the same number of images per class
- In the Custom_dataset(), **it was not ensured that the validation and training set have the same number of images per class**

How can this
problem be solved?


```


1  class Custom_Dataset_Balanced_Split(Dataset):
2      def __init__(self, split="train"):
3          if split == "train" or split == "val":
4              is_train = True
5
6
7          transform = transforms.ToTensor()
8          self.dataset = datasets.CIFAR10('data', train=is_train,
9                                         download=True, transform=transform)
10
11         # self.dataset.targets: returns the image classes (0,1,...,9)
12         # The classes are not ordered -> self.dataset.targets : [6,2,5,9,6,1,...,3]
13         x = np.array(self.dataset.targets)
14         # np.argsort(x) : return the indices ordered such that is possible to
15         # retrieve the classes like : [0,0,0,...,1,1,1,...,2,2,2,...,9,9,9]
16         sorted_indices = np.argsort(x)
17         num_elem_class = int(np.sum(x==0))
18         s = np.zeros((10,num_elem_class),int)
19         # For each class, store the corresponding indices in an array.
20         for i in range(10):
21             s[i] = sorted_indices[i*num_elem_class:(i+1)*num_elem_class]
22
23

```



```
1  indicez_zip = zip(s[0], s[1], s[2],s[3],s[4], s[5],s[6],s[7],s[8],s[9])
2
3
4  self.sorted_indices = [item for indices in indicez_zip for item in indices]
5
6  # for idx in range(len()):
7  #     print(self.dataset[self.sorted_indices[idx]][1])
8  # Output: 0,1,2,3,4,5,6,7,8,9,0,1,2,...,9
9
10 # lets use 20% of the training dataset to validation
11 if split == "train":
12     self.sorted_indices = self.sorted_indices[0:int(len(self.sorted_indices)*0.8)]
13 elif split == "val":
14     self.sorted_indices = self.sorted_indices[int(len(self.sorted_indices)*0.8):]
15
16 def __getitem__(self, idx):
17     image = self.dataset[self.sorted_indices[idx]][0]
18     label = self.dataset[self.sorted_indices[idx]][1]
19     return image, label
20
21 def __len__(self):
22     return len(self.sorted_indices)
23
24
```

2D Convolution



CLASS `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)` [SOURCE]

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{out}, H_{out}, W_{out})$ can be precisely described as:

$$\text{out}(N_i, C_{out_j}) = \text{bias}(C_{out_j}) + \sum_{k=0}^{C_{in}-1} \text{weight}(C_{out_j}, k) \star \text{input}(N_i, k)$$

- Input: $(N, C_{in}, H_{in}, W_{in})$ or (C_{in}, H_{in}, W_{in})
- Output: $(N, C_{out}, H_{out}, W_{out})$ or $(C_{out}, H_{out}, W_{out})$, where

$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 \times \text{padding}[1] - \text{dilation}[1] \times (\text{kernel_size}[1] - 1) - 1}{\text{stride}[1]} + 1 \right\rfloor$$

2D Convolution: Parameters

- **in_channels**: Number of channels in the input
- **out_channels**: Number of output channels
- **kernel_size**: Size of the convolving kernel
- **bias**: If True, adds a learnable bias to the output
- **dilation**: Spacing between kernel elements
- **padding**: Padding added to all four sides of the input.

Linear layer

```
CLASS torch.nn.Linear(in_features, out_features, bias=True, device=None, dtype=None) [SOURCE]
```

Applies an affine linear transformation to the incoming data: $y = xA^T + b$.

This module supports **TensorFloat32**.

On certain ROCm devices, when using float16 inputs this module will use **different precision** for backward.

Parameters

- **in_features** (*int*) – size of each input sample
- **out_features** (*int*) – size of each output sample
- **bias** (*bool*) – If set to `False`, the layer will not learn an additive bias. Default: `True`

Pooling layer

Max Pooling

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2
pool size

100	184
12	45


Average Pooling

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2
pool size


36	80
12	15

Activate functions




Sigmoid

$$g(z) = \frac{1}{1 + e^{-z}}$$




ReLU

$$g(z) = \max(0, z)$$



Tanh

$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$




Leaky ReLU

$$g(z) = \max(\epsilon z, z)$$

with $\epsilon \ll 1$

Model



```

1  class Net(nn.Module):
2      def __init__(self):
3          super().__init__()
4          self.conv1 = nn.Conv2d(3, 6, 5)
5          self.pool = nn.MaxPool2d(2, 2)
6          self.conv2 = nn.Conv2d(6, 16, 5)
7          # 5*5 : Resolution on the input in the first
8          # linear layer.
9          self.fc1 = nn.Linear(16 * 5 * 5, 120)
10         self.fc2 = nn.Linear(120, 84)
11         self.fc3 = nn.Linear(84, 10)
12
13     def forward(self, x):
14         x = self.pool(F.relu(self.conv1(x)))
15         x = self.pool(F.relu(self.conv2(x)))
16         # flatten all dimensions except batch
17         x = torch.flatten(x, 1)
18         x = F.relu(self.fc1(x))
19         x = F.relu(self.fc2(x))
20         x = self.fc3(x)
21         return x
22

```

Model: What if I want to add more layers to the Model?

- Create a new model that inherits the Net() class
- In the `__init__()` define the new additional layers
- Change the `forward()`

```
1 class Net_2(Net):
2     def __init__(self):
3         super(Net_2, self).__init__()
4         self.fc3 = nn.Linear(84, 42)
5         self.fc4 = nn.Linear(42, 10)
6
7     def forward(self, x):
8         x = self.pool((F.relu(self.conv1(x))))
9         x = self.pool((F.relu(self.conv2(x))))
10        x = torch.flatten(x, 1)
11        x = F.relu(self.fc1(x))
12        x = F.relu(self.fc2(x))
13        x = F.relu(self.fc3(x))
14        x = self.fc4(x)
15        return x
16
```



Let's learn how to train a classifier



What does the code for train a model looks like?

- Create dataset ✓

```
1 train_set = Custom_Dataset_Balanced_Split(split="train")
2 val_set = Custom_Dataset_Balanced_Split(split="val")
3 test_set = Custom_Dataset_Balanced_Split(split="test")
4
5 trainloader = DataLoader(train_set, batch_size=10,
6                           shuffle=False, num_workers=2)
7
8 valloader = DataLoader(val_set, batch_size=10,
9                        shuffle=False, num_workers=2)
10
11 testloader = DataLoader(test_set, batch_size=10,
12                          shuffle=False, num_workers=2)
```

What does the code for train a model looks like?

- Create dataset ✓
- Create a model ✓

```
1 device = "cuda" if torch.cuda.is_available() else "cpu"  
2  
3 net = Net()  
4  
5 net.to(device)
```

What does the code for train a model looks like?

- Create dataset ✓
- Create a model ✓
- Create optimiser

- PyTorch module that performs automatic differentiation on tensors.
- Keeps track of the operations performed on tensors and builds a computation graph
- It uses the the computation graph to calculate the gradient with respect to its inputs, allowing for backpropagation to be performed in deep learning models.

What does the code for train a model looks like?

- Create dataset ✓
- Create a model ✓
- Create optimiser ✓
- Define the loss ✓

```
1 optimizer = optim.Adam(net.parameters(), lr=0.001)
2
3 criterion = nn.CrossEntropyLoss()
```

What does the code for train a model looks like?

- Create dataset ✓
- Create a model ✓
- Create optimiser ✓
- Define the loss ✓
- Training Loop

Training Loop

```
1
2  for epoch in range(num_epochs): # loop over the dataset multiple times
3
4      running_loss = 0.0
5      correct = 0
6      total = 0
7      for i, data in enumerate(trainloader):
8          # get the inputs; data is a list of [inputs, labels]
9          inputs, labels = data
10         # Transfer the inputs to the device in use
11         inputs=inputs.to(device)
12         labels=labels.to(device)
13
14         # zero the parameter gradients
15         optimizer.zero_grad() # DO NOT FORGET!!
```

Training Loop

```

1      # forward + backward + optimize
2      outputs = net(inputs)
3      loss = criterion(outputs, labels)
4      loss.backward()
5      optimizer.step() # update model learnable weights
6
7
8      total += labels.size(0)
9      correct += accuracy(outputs)
10     # print statistics
11     running_loss += loss.item()
12     if i % 2000 == 1999:    # print every 2000 mini-batches
13         print(f'[{epoch + 1}, {i + 1:5d}] loss: {running_loss / 2000:.3f}')
14         running_loss = 0.0
15
16     print(f'Accuracy of the network on training set: {100 * correct / total} %')
```

What does the code for train a model looks like?

- Create dataset ✓
- Create a model ✓
- Create optimiser ✓
- Define the loss ✓
- Training Loop ✓
- Validation Loop ✓

Validation Loop

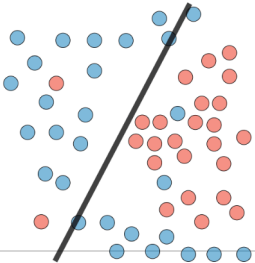
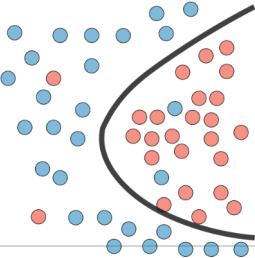
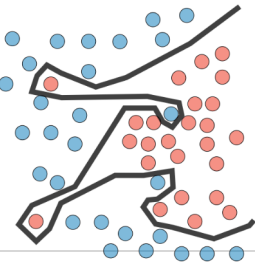
```
1
2 correct = 0
3 total = 0
4 # since we're not training, we don't need to calculate the gradients for our
  outputs
5 with torch.no_grad():
6     for data in valloader:
7
8         images, labels = data
9         images=images.to(device)
10        labels=labels.to(device)
11        # calculate outputs by running images through the network
12        outputs = net(images)
13
14        total += labels.size(0)
15        correct += accuracy(outputs)
16
17 accuracy_val = correct / total
```



Validation Loop: Save checkpoint

```
1  # use the validation
2  if accuracy_val > accuracy_val_min:
3      accuracy_val_min = accuracy_val
4      # Save the model weights
5      # The best checkpoint will be later
6      # used for inference
7      torch.save(net.state_dict(), checkpoint_path)
8      best_epoch = epoch
9
10 print(f'Accuracy of the network on validation set: {np.round(100 * accuracy_val,3)}
    %')
```

Model fitting: Classification of red vs blue dots

Underfitting	Just right	Overfitting
<ul style="list-style-type: none">• High training error• Training error close to validation error• High bias	<ul style="list-style-type: none">• Training error slightly lower than validation error	<ul style="list-style-type: none">• Very low training error• Training error much lower than validation error
		

Model fitting (what we generally see)



What does the code for the inference look like?

- Create dataset ✓
- Create a model ✓
- Load checkpoint

```
1 net.load_state_dict(torch.load(checkpoint_path))  
2 classes = ('plane', 'car', 'bird', 'cat',  
3           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

What does the code for the inference look like?

- Create dataset ✓
- Create a model ✓
- Load checkpoint ✓
- Inference loop

```

1  correct = 0
2  total = 0
3  # since we're not training, we don't need to calculate
   the gradients for our outputs
4  with torch.no_grad():
5      for data in testloader:
6          images, labels = data
7          images=images.to(device)
8          labels=labels.to(device)
9          # calculate outputs
10         outputs = net(images)
11         # Class with the highest energy is what we
12         # choose as prediction
13         _, predicted = torch.max(outputs.data, 1)
14         total += labels.size(0)
15         correct += (predicted == labels).sum().item()
16
17     print(f'Accuracy of the network on the 10000 test images:
        {100 * correct / total} %')
```

What does the code for the inference look like?

- Create dataset ✓
- Create a model ✓
- Load checkpoint ✓
- Inference loop ✓
- Performance per class

```

1  # prepare to count predictions for each class
2  correct_pred = {classname: 0 for classname in classes}
3  total_pred = {classname: 0 for classname in classes}
4
5  # again no gradients needed
6  with torch.no_grad():
7      for data in testloader:
8          images, labels = data
9          images=images.to(device)
10         labels=labels.to(device)
11         outputs = net(images)
12         _, predictions = torch.max(outputs, 1)
13         # collect the correct predictions for each class
14         for label, prediction in zip(labels, predictions)
15             :
16                 if label == prediction:
17                     correct_pred[classes[label]] += 1
18                     total_pred[classes[label]] += 1

```

What does the code for the inference look like?

- Create dataset ✓
- Create a model ✓
- Load checkpoint ✓
- Inference loop ✓
- Performance per class ✓

```
1 # print accuracy for each class
2 for classname, correct_count in correct_pred.items():
3     accuracy = 100 * float(correct_count) / total_pred[
4         classname]
5     print(f'Accuracy for class: {classname:5s} is {
6         accuracy:.1f} %')
```

Thank You